

[shared externally]

- All state are defined by
https://sourcegraph.com/github.com/ray-project/ray@06ef4ab94e09738c4afe4d4fdc75a9cb41464b4c/-/blob/python/ray/workflow/workflow_state.py?L49:7
- Submitting a workflow
 - Submit the DAG to WMA(WorkflowManagementActor found in workflow_access)
 - WMA -> write it into DB
 - WMA -> create a state and run
https://sourcegraph.com/github.com/ray-project/ray@HEAD/-/blob/python/ray/workflow/workflow_executor.py
 -
- Dashboard (<https://github.com/ray-project/ray/tree/master/dashboard>)
 - Running in the head node
 - No driver here!
- Dashboard Agent (<https://github.com/ray-project/ray/tree/master/dashboard>)
 - Running in the work node (head node).
 - Each raylet will have one
 - Ray.init can run here
 - Agent can have a driver
- For dashboard and the dashboard agent, you can check this:
 - <https://github.com/ray-project/ray/tree/master/dashboard/modules/log>
 - Dashboard Agent
 - https://github.com/ray-project/ray/blob/master/dashboard/modules/log/log_agent.py
 - Log agent doing some work for logs
 - Dashboard
 - https://github.com/ray-project/ray/blob/master/dashboard/modules/log/log_head.py
 - Log head pulling the request from log agent
 - Serve Agent
 - https://github.com/ray-project/ray/blob/master/dashboard/modules/serve/serve_agent.py#L168
 - Here we access serve controller actor handle explicitly
 - You probably want to do the similar things for workflow actor manager
 - Call this to get actor handle:
 - https://sourcegraph.com/github.com/ray-project/ray@06ef4ab94e09738c4afe4d4fdc75a9cb41464b4c/-/blob/python/ray/workflow/workflow_access.py?L371-374
- Dashboard —> some request —> Dashboard Agent (ray.init) – send request to WorkflowManagementActor —> read the data from the memory for running workflow (read from storage for not running one [second step])
- Dashboard Agent
 - Define a WorkflowObservability module

- `func(workflow_id) -> (dict(task_id, metadata))`
- Inside the module, you'll have a workflow management actor handler.
- You need to define the methods in WMA to retrieve the state of a workflow job
- Then you'll send the useful state back to the dashboard

MVPs

- P0
 - Visualize the running workflow DAG
 - Check the progress of the running workflow
 - Some metadata of the workflow tasks.
 - Dynamic updates?
 - In dashboard agent, you can send a remote call to WMA (request new data)
 - Return if the state of some workflow changed
- P1
 - Check the status of the non-running DAG
 - Failed? Get the exception or error message?
 - Check the metadata of the workflow tasks
 - How long does it run
 - The size of the output
- P2
 - Workflow management
 - Kill a workflow
 - Restart/resume